

Bayesian Data Analysis Chapter 2

January 19, 2019

```
In [1]: %matplotlib inline
import pymc3 as pm
import numpy as np
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib
palette = 'muted'
sns.set_palette(palette); sns.set_color_codes(palette)

figureSize=(16,10)

In [2]: def posterior_grid(grid_points=100, heads=6, tosses=9):
    """
    A grid implementation for the coin-flip problem
    """
    # define a grid
    grid = np.linspace(0, 1, grid_points)

    # define prior
    prior = np.repeat(5, grid_points) # uniform
    #prior = (grid <= 0.4).astype(int) # truncated
    #prior = abs(grid - 0.5) # "M" prior

    # compute likelihood at each point in the grid
    likelihood = stats.binom.pmf(heads, tosses, grid)

    # compute product of likelihood and prior
    unstd_posterior = likelihood * prior


    # standardize the posterior, so it sums to 1
    posterior = unstd_posterior / unstd_posterior.sum()
    return grid, posterior

In [3]: points = 15
h, n = 1, 4
grid, posterior = posterior_grid(points, h, n)
```

```

plt.plot(grid, posterior, 'o-')
plt.plot(0, 0, label='heads = {} \ntosses = {}'.format(h, n), alpha=0)
plt.xlabel(r'$\theta$', fontsize=14)
plt.legend(loc=0, fontsize=14)
fig = matplotlib.pyplot.gcf()
fig.set_size_inches(figsize[0],figsize[1], forward=True)

```



0.1 Monte Carlo

In [4]: N = 10000


```

x, y = np.random.uniform(-1, 1, size=(2, N))
inside = (x**2 + y**2) <= 1
pi = inside.sum()*4/N
error = abs((pi - np.pi)/pi)* 100

outside = np.invert(inside)

plt.plot(x[inside], y[inside], 'b.')
plt.plot(x[outside], y[outside], 'r.')
plt.plot(0, 0, label='$\hat{\pi}$ = {:.3f} \nerror = {:.3f}%'.format(pi, error), alpha=0)
plt.axis('square')
plt.legend(frameon=True, framealpha=0.9, fontsize=16);
fig = matplotlib.pyplot.gcf()
fig.set_size_inches(figsize[0],figsize[1])

```



0.2 Metropolis

0.2.1 this is not a good implementation see <https://www.youtube.com/watch?v=OTO1DygELpY>

```
In [5]: def metropolis(func, steps=10000):  
        """A very simple Metropolis implementation"""  
        samples = np.zeros(steps)  
        old_x = func.mean()  
        old_prob = func.pdf(old_x)  
  
        for i in range(steps):  
            new_x = old_x + np.random.normal(0, 1)  
            new_prob = func.pdf(new_x)  
            acceptance = new_prob/old_prob
```

```

        if acceptance >= np.random.random():
            samples[i] = new_x
            old_x = new_x
            old_prob = new_prob
        else:
            samples[i] = old_x
    return samples


```

```

In [6]: np.random.seed(345)
        func = stats.beta(0.4, 2)
        samples = metropolis(func=func)
        x = np.linspace(0.01, .99, 100)
        y = func.pdf(x)
        plt.xlim(0, 1)
        plt.plot(x, y, 'r-', lw=3, label='True distribution')
        plt.hist(samples, bins=30, normed=True, label='Estimated distribution')
        plt.xlabel('$x$', fontsize=14)
        plt.ylabel('$pdf(x)$', fontsize=14)
        plt.legend(fontsize=14)
        fig = matplotlib.pyplot.gcf()
        fig.set_size_inches(figsize[0],figsize[1])

```

/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/matplotlib/axes/_axes.py:6521
The 'normed' kwarg was deprecated in Matplotlib 2.1 and will be removed in 3.1. Use 'density' :
alternative="density", removal="3.1")




0.3 PyMC3 primer

```
In [7]: np.random.seed(123)
        n_experiments = 4
        theta_real = 0.35 # unkwon value in a real experiment
        data = stats.bernoulli.rvs(p=theta_real, size=n_experiments)
```

```
In [8]: with pm.Model() as our_first_model:
        # a priori
        theta = pm.Beta('theta', alpha=1, beta=1)
        # likelihood
        y = pm.Bernoulli('y', p=theta, observed=data)
        # start at the maximum a posteriori probability (MAP) estimate
        # see https://en.wikipedia.org/wiki/Maximum_a_posteriori_estimation
        start = pm.find_MAP()
        # use M-H algorithm to sample
        step = pm.Metropolis()
        trace = pm.sample(1000, step=step, start=start)
```

```
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/pymc3/tuning/starting.py:61: UserWarning: find_MAP should not be used to initialize the NUTS sampler, simply call pymc3
logp = -2.7726, ||grad|| = 1: 100%|| 6/6 [00:00<00:00, 3584.36it/s]
Multiprocess sampling (4 chains in 4 jobs)
Metropolis: [theta]
Sampling 4 chains: 100%|| 6000/6000 [00:00<00:00, 12985.97draws/s]
The number of effective samples is smaller than 25% for some parameters.
```


```
In [9]: burnin = 0 # no burnin
        chain = trace[burnin:]
        pm.traceplot(chain, lines={'theta':theta_real});
```



```
In [10]: with our_first_model:
         step = pm.Metropolis()
         multi_trace = pm.sample(10000, step=step, njobs=8)
```

```
Multiprocess sampling (8 chains in 8 jobs)
Metropolis: [theta]
Sampling 8 chains: 100%|| 84000/84000 [00:06<00:00, 12045.26draws/s]
The number of effective samples is smaller than 25% for some parameters.
```


```
In [11]: burnin = 0 # no burnin
multi_chain = multi_trace[burnin:]
pm.traceplot(multi_chain, lines={'theta':theta_real});
```



```
In [12]: pm.gelman_rubin(multi_chain)
```

```
Out[12]: {'theta': 1.000235781648356}
```

```
In [13]: pm.forestplot(multi_chain, varnames=['theta']);
```



```
In [14]: pm.summary(multi_chain)
```

```
Out[14]:
```

	mean	sd	mc_error	hpd_2.5	hpd_97.5	n_eff	Rhat
theta	0.335045	0.17943	0.001399	0.023004	0.668089	16414.812056	1.000236


```
In [15]: pm.summary(multi_chain)
```

```
Out[15]:
```

	mean	sd	mc_error	hpd_2.5	hpd_97.5	n_eff	Rhat
theta	0.335045	0.17943	0.001399	0.023004	0.668089	16414.812056	1.000236

```
In [16]: pm.autocorrplot(chain)
```

```
Out[16]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f3b21839c50>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x7f3b2171d710>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x7f3b216486a0>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x7f3b21671630>]],  
              dtype=object)
```

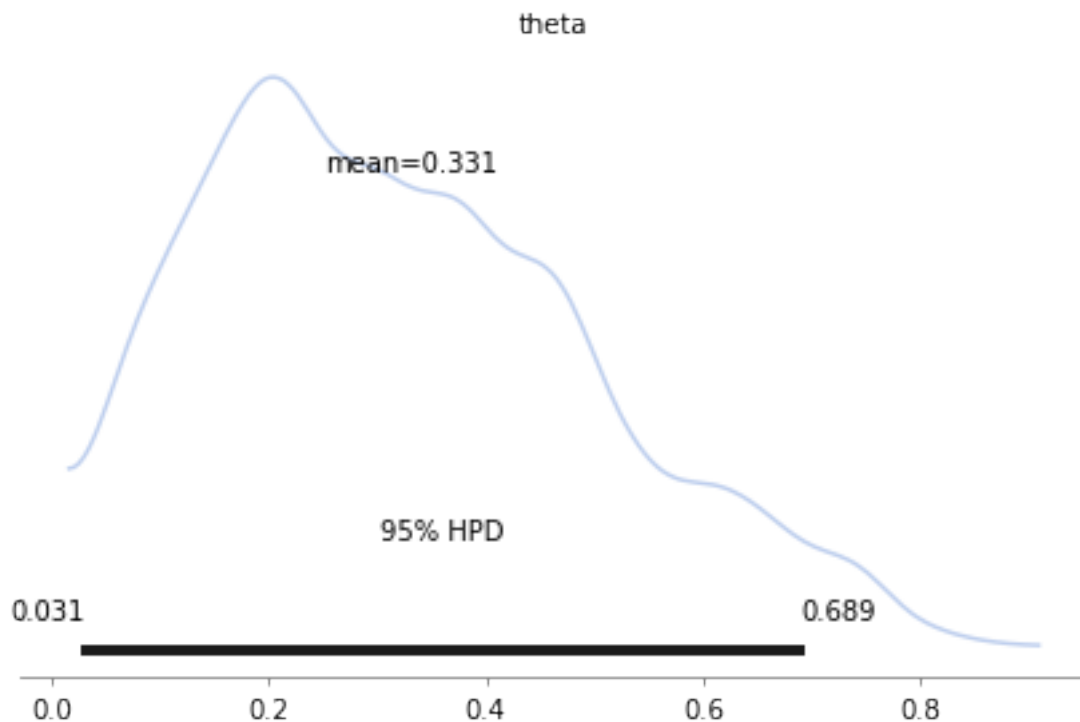


```
In [17]: pm.effective_n(multi_chain)['theta']
```

```
Out[17]: 16414.812055918923
```

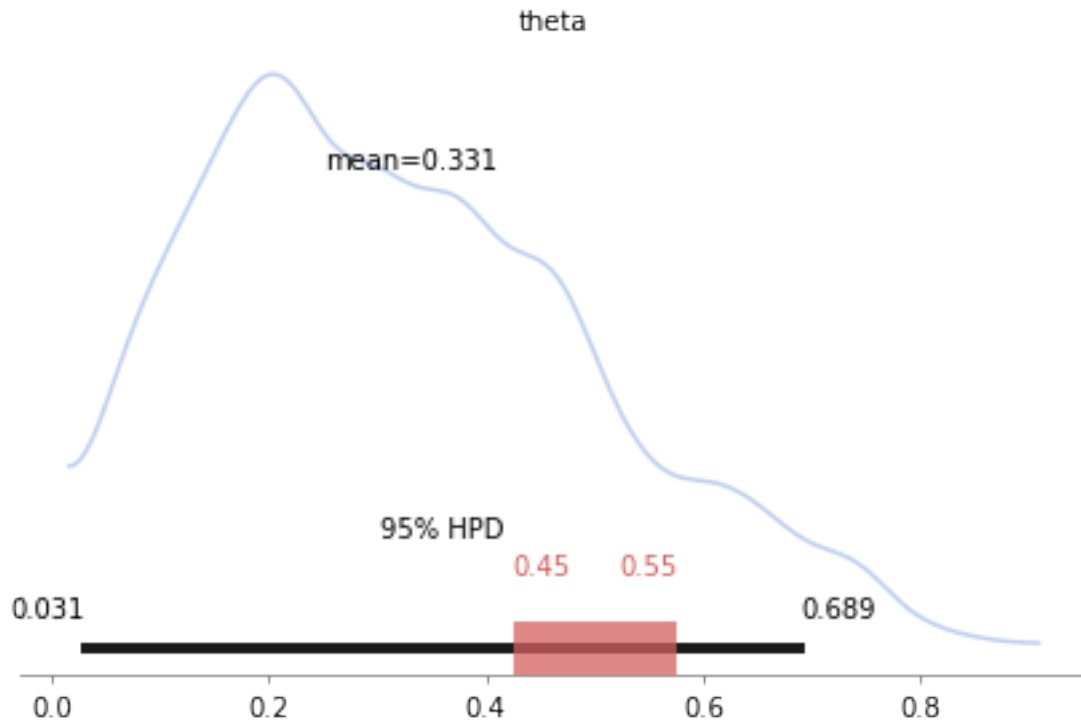
```
In [18]: pm.plot_posterior(chain, kde_plot=True)
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3b218d6d68>
```




```
In [19]: pm.plot_posterior(chain, kde_plot=True, rope=[0.45, .55])
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3b229dd0f0>
```



```
In [20]: pm.plot_posterior(chain, kde_plot=True, ref_val=0.5)
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3b22acdc50>
```

```
In [21]: import sys, IPython, scipy, matplotlib, platform
         print("This notebook was created on a %s computer running %s and using:\nPython %s\nIP
```

This notebook was created on a x86_64 computer running debian buster/sid and using:

Python 3.7.2

IPython 7.2.0

PyMC3 3.6

NumPy 1.16.0

SciPy 1.2.0

Matplotlib 3.0.2

Seaborn 0.9.0

/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/ipykernel_launcher.py:2: DeprecationWarning: The 'warn' argument is deprecated since version 3.0. Use 'warn_kwarg='warn_kwarg' instead.