


Bayesian Data Analysis Chapter 5

January 20, 2019

1 Logistic Regression

```
In [1]: %matplotlib inline
import pymc3 as pm
import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns
import theano.tensor as tt
palette = 'muted'
sns.set_palette(palette); sns.set_color_codes(palette)
np.set_printoptions(precision=2)
pd.set_option('display.precision', 2)

In [2]: z = np.linspace(-10, 10, 100)
logistic = 1 / (1 + np.exp(-z))
plt.plot(z, logistic);
plt.xlabel('$z$', fontsize=16);
plt.ylabel('$\text{logistic}(z)$', fontsize=16);
```




```
In [3]: iris = sns.load_dataset("iris")
iris.head()
```

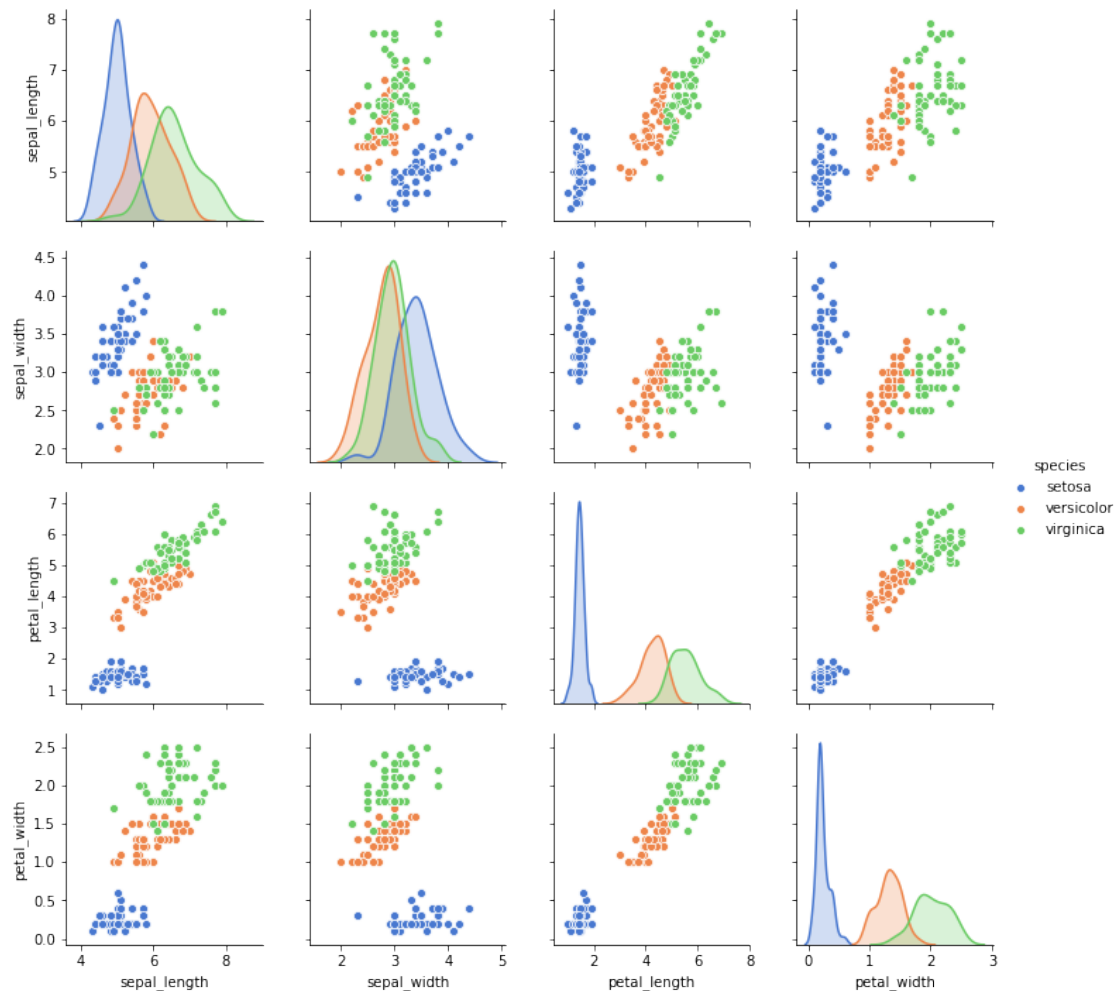
```
Out[3]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
In [4]: sns.stripplot(x="species", y="sepal_length", data=iris, jitter=True);
```



```
In [5]: sns.pairplot(iris, hue='species', diag_kind='kde');
```



```
In [6]: df = iris.query("species == ('setosa', 'versicolor')")
        y_0 = pd.Categorical(df['species']).codes
        x_n = 'sepal_length'
        x_0 = df[x_n].values
```

```
In [7]: with pm.Model() as model_0:
        alpha = pm.Normal('alpha', mu=0, sd=10)
        beta = pm.Normal('beta', mu=0, sd=10)

        mu = alpha + pm.math.dot(x_0, beta)
        #theta = pm.Deterministic('theta', tt.nnet.sigmoid(mu))
        # the theta of the Bernoulli is derived from the logistic inverse link function
        theta = pm.Deterministic('theta', 1 / (1 + pm.math.exp(-mu)))

        bd = pm.Deterministic('bd', -alpha/beta)
```

```

yl = pm.Bernoulli('y1', p=theta, observed=y_0)

#     start = pm.find_MAP()
step = pm.NUTS()
trace_0 = pm.sample(5000, step, chains=1, njobs=1)

```

Sequential sampling (1 chains in 1 job)

NUTS: [beta, alpha]

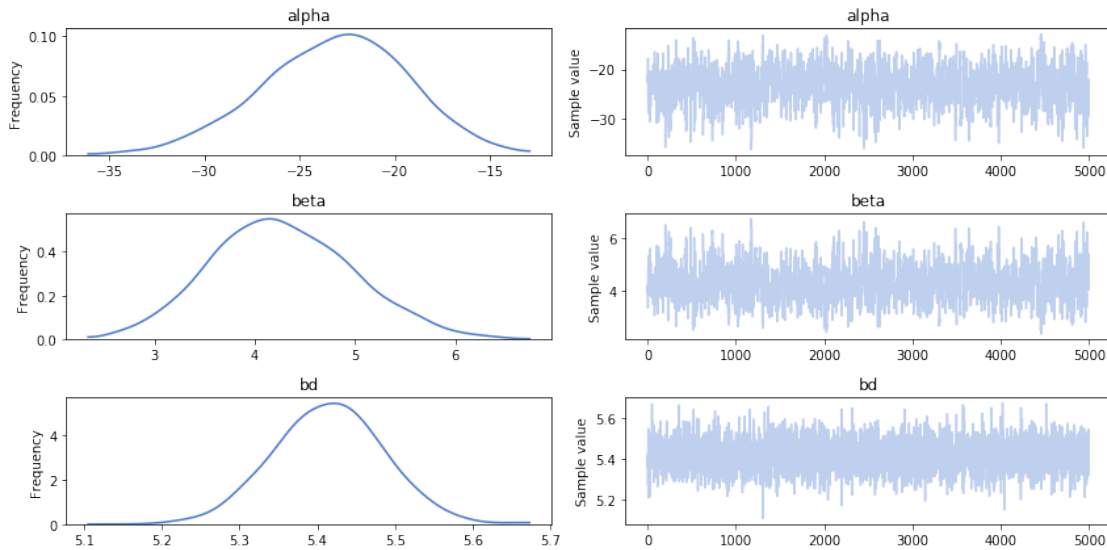
100%|| 5500/5500 [00:14<00:00, 367.59it/s]

Only one chain was sampled, this makes it impossible to run some convergence checks

```

In [8]: varnames = ['alpha', 'beta', 'bd']
pm.traceplot(trace_0, varnames);

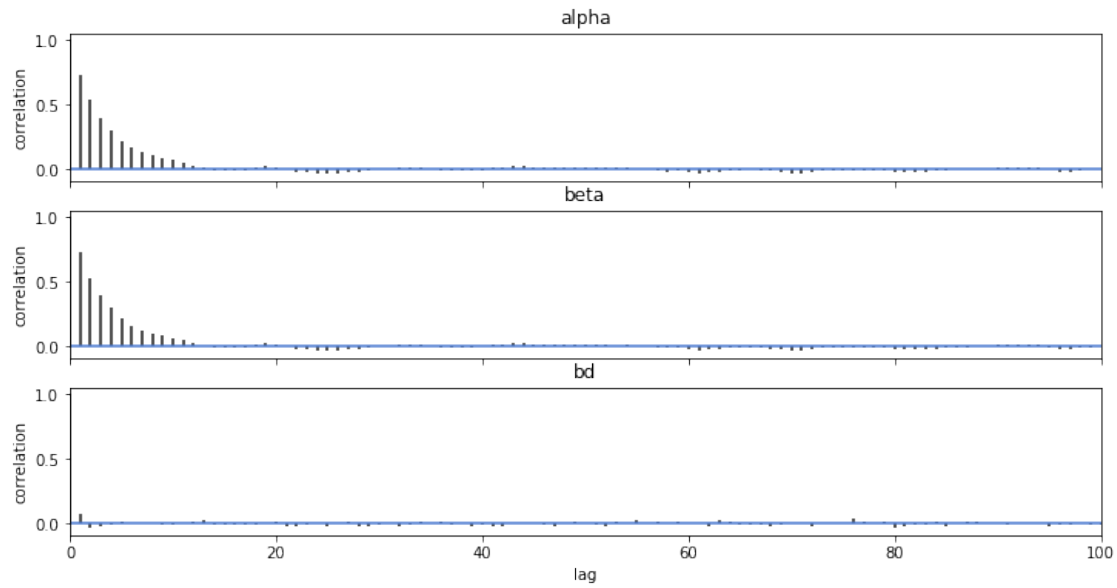
```




```

In [9]: pm.autocorrplot(trace_0, varnames);

```



```
In [10]: sns.kdeplot(trace_0['alpha'], trace_0['beta']);
```



```
In [11]: pm.summary(trace_0, varnames)
```

```
Out[11]:
```


	mean	sd	mc_error	hpd_2.5	hpd_97.5
alpha	-23.17	3.84	1.31e-01	-30.82	-15.89

beta	4.28	0.71	2.42e-02	2.94	5.70
bd	5.42	0.07	9.42e-04	5.28	5.55

```
In [12]: theta = trace_0['theta'].mean(axis=0)
idx = np.argsort(x_0)
plt.plot(x_0[idx], theta[idx], color='b', lw=3);
plt.axvline(trace_0['bd'].mean(), ymax=1, color='r');
bd_hpd = pm.hpd(trace_0['bd'])
plt.fill_betweenx([0, 1], bd_hpd[0], bd_hpd[1], color='r', alpha=0.5);

plt.plot(x_0, y_0, 'o', color='k');
theta_hpd = pm.hpd(trace_0['theta'])[idx]
plt.fill_between(x_0[idx], theta_hpd[:,0], theta_hpd[:,1], color='b', alpha=0.5);

plt.xlabel(x_n, fontsize=16);
plt.ylabel(r'$\theta$', rotation=0, fontsize=16);
```



```
In [13]: def classify(n, threshold):
        """
        A simple classifying function
        """
        n = np.array(n)
        mu = trace_0['alpha'].mean() + trace_0['beta'].mean() * n
        prob = 1 / (1 + np.exp(-mu))
```

```

        return prob, prob >= threshold

classify([5, 5.5, 6], 0.5)

Out[13]: (array([0.15, 0.59, 0.92]), array([False,  True,  True]))

```

1.1 Multiple logistic regression

```

In [14]: df = iris.query("species == ('setosa', 'versicolor')")
        y_1 = pd.Categorical(df['species']).codes
        x_n = ['sepal_length', 'sepal_width'] #['petal_length', 'petal_width']
        x_1 = df[x_n].values

        #x_1 = (x_1 - x_1.mean(axis=0))/x_1.std(axis=0) #
        #x_1 = (x_1 - x_1.mean(axis=0))

In [15]: with pm.Model() as model_1:
        # We define the priors
        alpha = pm.Normal('alpha', mu=0, sd=10)
        beta = pm.Normal('beta', mu=0, sd=2, shape=len(x_n))

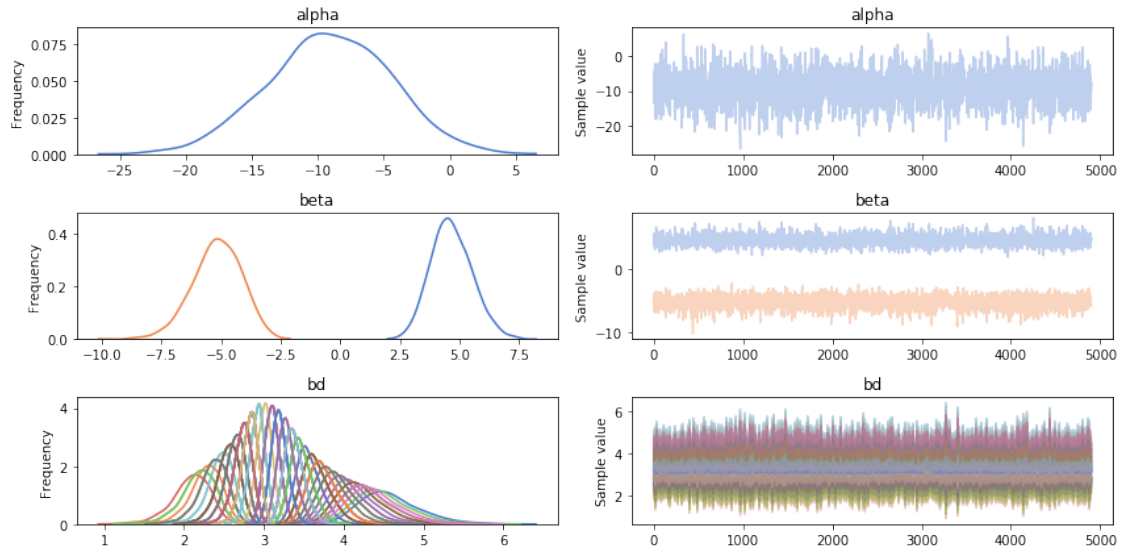
        mu = alpha + pm.math.dot(x_1, beta)
        # Apply the logistic linking function
        theta = 1 / (1 + pm.math.exp(-mu))
        # Compute the boundary decision
        bd = pm.Deterministic('bd', -alpha/beta[1] - beta[0]/beta[1] * x_1[:,0])

        # Define the likelihood
        y1 = pm.Bernoulli('y1', p=theta, observed=y_1)
        # Sampling
        trace_1 = pm.sample(5000,chains=1, njobs=1)

Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Sequential sampling (1 chains in 1 job)
NUTS: [beta, alpha]
100%|| 5500/5500 [00:20<00:00, 272.09it/s]
The acceptance probability does not match the target. It is 0.8931604002334997, but should be 0.8
Only one chain was sampled, this makes it impossible to run some convergence checks

In [16]: chain_1 = trace_1[100:]
        pm.traceplot(chain_1);

```

Note that now the boundary decision has 100 parameters, one for each data point

```
In [17]: pm.summary(chain_1, ['alpha', 'beta'])
```

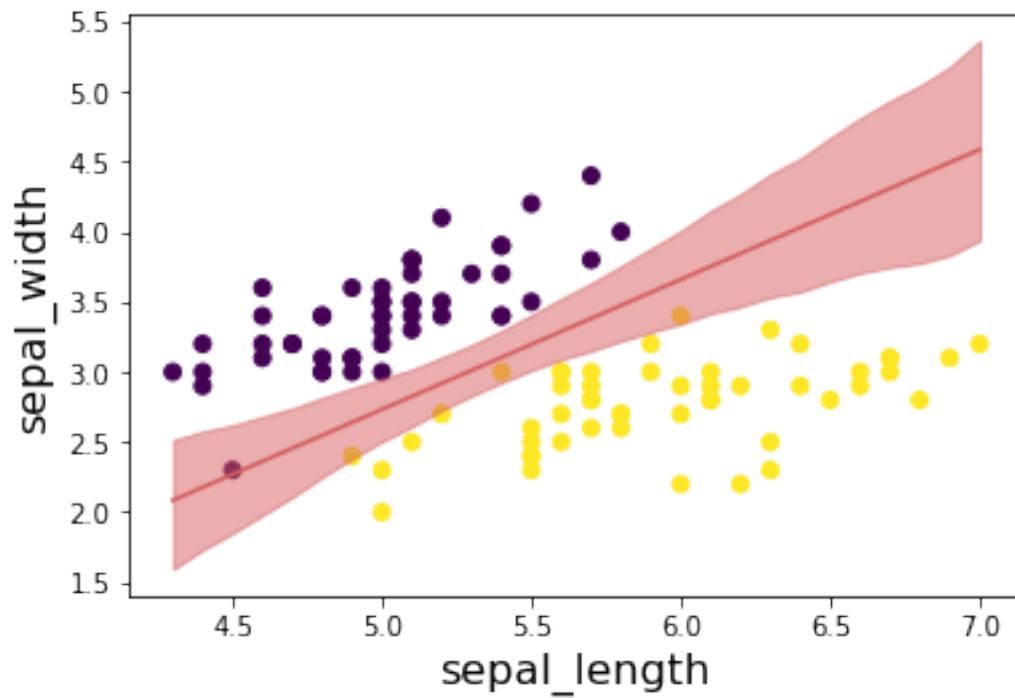
```
Out[17]:
```

	mean	sd	mc_error	hpd_2.5	hpd_97.5
alpha	-9.12	4.67	0.11	-18.29	-0.17
beta__0	4.67	0.86	0.02	3.06	6.39
beta__1	-5.18	1.03	0.02	-7.18	-3.20

```
In [18]: idx = np.argsort(x_1[:,0])
         ld = chain_1['bd'].mean(0)[idx]
         plt.scatter(x_1[:,0], x_1[:,1], c=y_0);
         plt.plot(x_1[:,0][idx], ld, color='r');

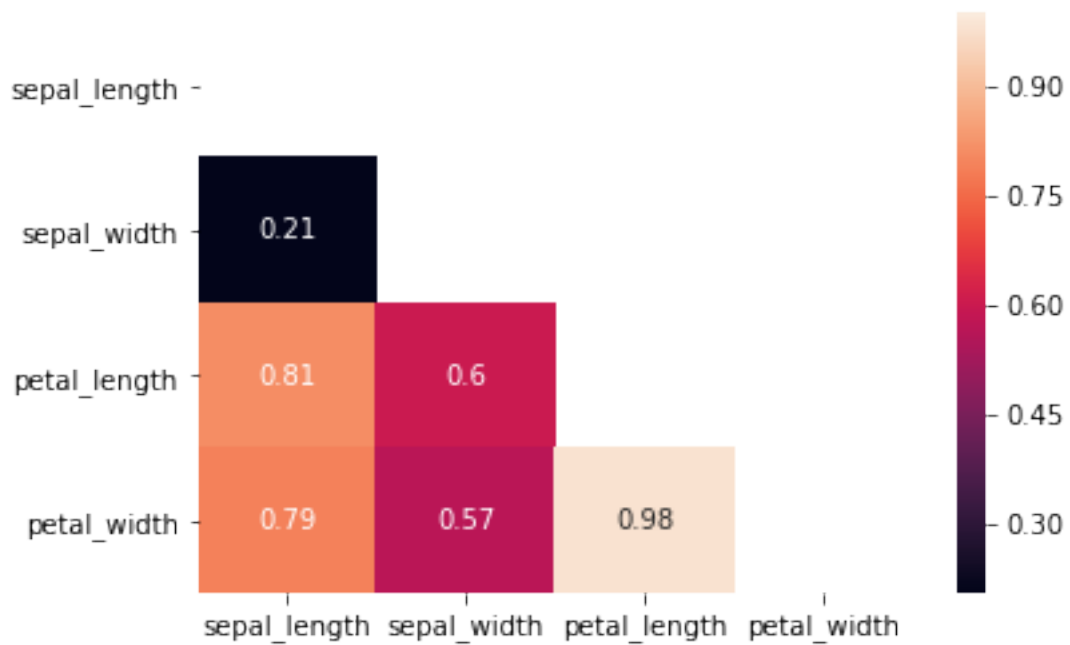
         ld_hpd = pm.hpd(chain_1['bd'])[idx]
         plt.fill_between(x_1[:,0][idx], ld_hpd[:,0], ld_hpd[:,1], color='r', alpha=0.5);

         plt.xlabel(x_n[0], fontsize=16);
         plt.ylabel(x_n[1], fontsize=16);
```



1.2 Correlated variables

```
In [19]: corr = iris[iris['species'] != 'virginica'].corr()
          mask = np.tri(*corr.shape).T
          sns.heatmap(corr.abs(), mask=mask, annot=True);
```



1.3 Unbalanced Classes

```
In [20]: df = iris.query("species == ('setosa', 'versicolor')")
df = df[45:] #df[22:78]
y_3 = pd.Categorical(df['species']).codes
x_n = ['sepal_length', 'sepal_width']
x_3 = df[x_n].values

In [21]: with pm.Model() as model_3:
    # We define the a prioris
    alpha = pm.Normal('alpha', mu=0, sd=10)
    beta = pm.Normal('beta', mu=0, sd=2, shape=len(x_n))

    mu = alpha + pm.math.dot(x_3, beta)
    # Aplly the logistic linking function
    p = 1 / (1 + pm.math.exp(-mu))
    # Compute the boundary decision
    ld = pm.Deterministic('ld', -alpha/beta[1] - beta[0]/beta[1] * x_3[:,0])

    # Define the likelihood
    y1 = pm.Bernoulli('y1', p=p, observed=y_3)
    trace_3 = pm.sample(5000,chains=1, njobs=1)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

Sequential sampling (1 chains in 1 job)

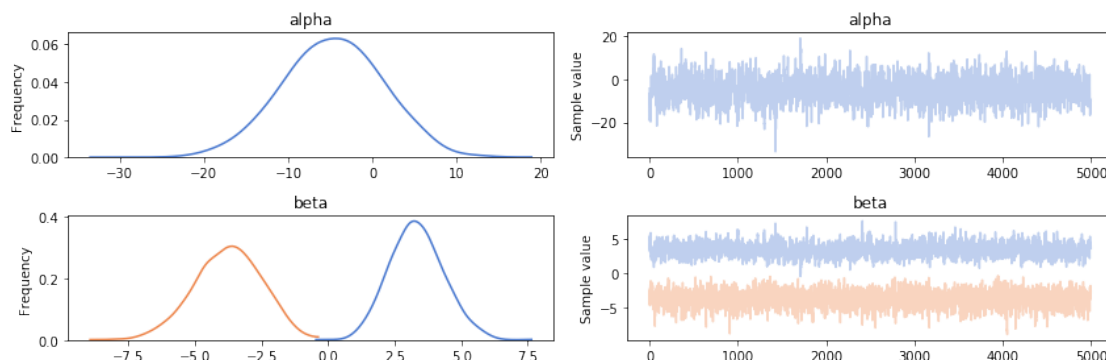
NUTS: [beta, alpha]

100%|| 5500/5500 [00:14<00:00, 383.43it/s]

There were 112 divergences after tuning. Increase `target_accept` or reparameterize.

Only one chain was sampled, this makes it impossible to run some convergence checks

```
In [22]: cadena_3 = trace_3[:]
varnames = ['alpha', 'beta']
pm.traceplot(cadena_3, varnames);
```



```
In [23]: pm.summary(cadena_3, varnames)
```


```
Out[23]:
```

	mean	sd	mc_error	hpd_2.5	hpd_97.5
alpha	-4.70	5.94	0.18	-15.93	6.74
beta__0	3.33	1.01	0.03	1.30	5.25
beta__1	-3.77	1.27	0.03	-6.22	-1.36

```
In [24]: idx = np.argsort(x_3[:,0])
ld = trace_3['ld'].mean(0)[idx]
plt.scatter(x_3[:,0], x_3[:,1], c=y_3);
plt.plot(x_3[:,0][idx], ld, color='r');

ld_hpd = pm.hpd(trace_3['ld'])[idx]
plt.fill_between(x_3[:,0][idx], ld_hpd[:,0], ld_hpd[:,1], color='r', alpha=0.5);

plt.xlabel(x_n[0], fontsize=16);
plt.ylabel(x_n[1], fontsize=16);
```



1.4 Softmax logistic regression (AKA multinomial regression)

```
In [25]: iris = sns.load_dataset("iris")
y_s = pd.Categorical(iris['species']).codes
```

```

x_n = iris.columns[:-1]
x_s = iris[x_n].values
x_s = (x_s - x_s.mean(axis=0))/x_s.std(axis=0)
#x_s = (x_s - x_s.mean(axis=0))

```

```

In [26]: with pm.Model() as model_s:
    alpha = pm.Normal('alpha', mu=0, sd=2, shape=3)
    beta = pm.Normal('beta', mu=0, sd=2, shape=(4,3))

    mu = alpha + pm.math.dot(x_s, beta)
    theta = tt.nnet.softmax(mu)

    y1 = pm.Categorical('y1', p=theta, observed=y_s)
    step = pm.NUTS()
    trace_s = pm.sample(2000, step, chains=1, njobs=1)

```

```

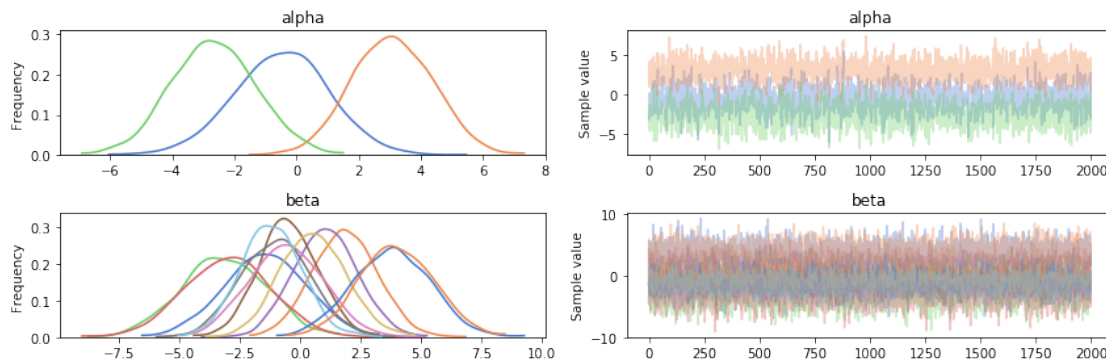
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/theano/tensor/subtensor.py:21
    rval = inputs[0].__getitem__(inputs[1:])
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/theano/tensor/subtensor.py:21
    rval = inputs[0].__getitem__(inputs[1:])
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/theano/tensor/subtensor.py:21
    rval = inputs[0].__getitem__(inputs[1:])
Sequential sampling (1 chains in 1 job)
NUTS: [beta, alpha]
100%|| 2500/2500 [00:12<00:00, 197.64it/s]
Only one chain was sampled, this makes it impossible to run some convergence checks

```

```

In [27]: pm.traceplot(trace_s);

```



```

In [28]: pm.summary(trace_s)

```


```

Out[28]:
          mean      sd  mc_error  hpd_2.5  hpd_97.5
alpha__0 -0.50  1.50     0.04    -3.64     2.20
alpha__1  3.07  1.29     0.03     0.56     5.51

```

alpha__2	-2.68	1.35	0.03	-5.19	0.08
beta__0_0	-1.45	1.68	0.04	-4.87	1.64
beta__0_1	1.04	1.29	0.04	-1.41	3.52
beta__0_2	0.53	1.34	0.04	-2.24	3.02
beta__1_0	1.85	1.31	0.03	-0.64	4.39
beta__1_1	-0.50	1.21	0.03	-2.77	1.95
beta__1_2	-1.20	1.25	0.03	-3.55	1.31
beta__2_0	-3.17	1.73	0.04	-6.72	-0.08
beta__2_1	-0.59	1.53	0.03	-3.63	2.25
beta__2_2	3.77	1.60	0.03	0.45	6.67
beta__3_0	-3.00	1.76	0.03	-6.57	0.26
beta__3_1	-1.02	1.42	0.03	-3.72	1.80
beta__3_2	3.98	1.52	0.03	1.23	7.22

```
In [29]: pm.autocorrplot(trace_s, varnames=['alpha']);
```



```
In [30]: probs = trace_s['alpha'].mean(axis=0) + np.dot(x_s, trace_s['beta'].mean(axis=0))
y_pred = []
for point in probs:
    # each array appended contains the softmax estimates of each class
    y_pred.append(np.exp(point)/np.sum(np.exp(point), axis=0))
print('Accuracy on training data is {0:.2%}'.format(np.sum(y_s == np.argmax(y_pred, axis=0))
```

Accuracy on training data is 97.33%

```
In [31]: with pm.Model() as model_sf:
    alpha = pm.Normal('alpha', mu=0, sd=2, shape=2)
```

```

beta = pm.Normal('beta', mu=0, sd=2, shape=(4,2))

alpha_f = tt.concatenate([[0] , alpha])
beta_f = tt.concatenate([np.zeros((4,1)) , beta], axis=1)

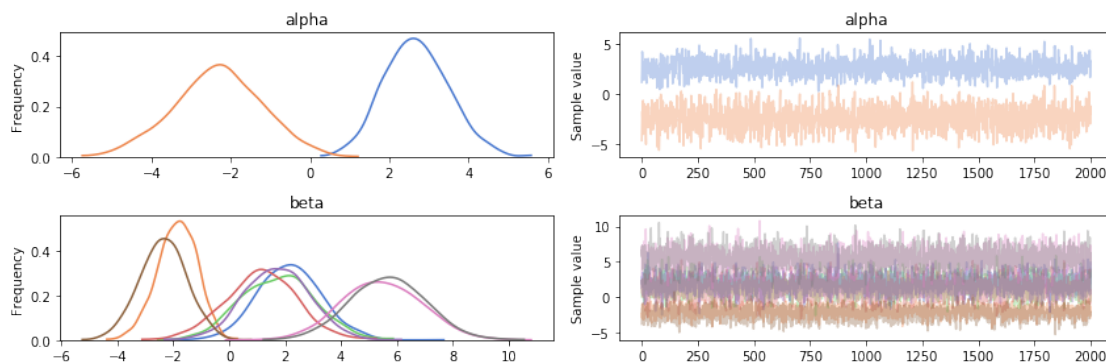
mu = alpha_f + pm.math.dot(x_s, beta_f)
theta = tt.nnet.softmax(mu)

y_l = pm.Categorical('y_l', p=theta, observed=y_s)
step = pm.NUTS()
trace_sf = pm.sample(2000, step, chains=1, njobs=1)

/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/theano/tensor/subtensor.py:21:
  rval = inputs[0].__getitem__(inputs[1:])
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/theano/tensor/subtensor.py:21:
  rval = inputs[0].__getitem__(inputs[1:])
Sequential sampling (1 chains in 1 job)
NUTS: [beta, alpha]
100%|| 2500/2500 [00:11<00:00, 222.00it/s]
The acceptance probability does not match the target. It is 0.8851232083982429, but should be 0.8
Only one chain was sampled, this makes it impossible to run some convergence checks

```

```
In [32]: pm.traceplot(trace_sf);
```



we get much better defined α 's and β 's than before because we didn't parametrize the problem correctly. We need to fix the number of parameters to $num_{classes} - 1$, since all probabilities must sum to one we had too many parameters in the first instance of the model. We do this by setting the first α and first β 's to zero (see code).

```
In [33]: pm.summary(trace_sf)
```

```

Out [33]:
          mean      sd  mc_error  hpd_2.5  hpd_97.5
alpha__0  2.68  0.82      0.02  1.12e+00  4.35

```

alpha__1	-2.34	1.12	0.03	-4.49e+00	-0.09
beta__0_0	2.14	1.12	0.03	2.37e-03	4.47
beta__0_1	1.69	1.16	0.03	-5.15e-01	4.04
beta__1_0	-1.88	0.70	0.02	-3.23e+00	-0.52
beta__1_1	-2.37	0.84	0.02	-3.99e+00	-0.71
beta__2_0	1.73	1.30	0.03	-6.13e-01	4.34
beta__2_1	5.43	1.42	0.04	2.49e+00	7.98
beta__3_0	1.12	1.27	0.03	-1.38e+00	3.65
beta__3_1	5.66	1.35	0.03	3.11e+00	8.40

1.5 Linear discriminat Analysis

```
In [34]: iris = sns.load_dataset("iris")
df = iris.query("species == ('setosa', 'versicolor')")
y_0 = pd.Categorical(df['species']).codes
x_n = 'sepal_length'
x_0 = df[x_n].values

In [35]: with pm.Model() as model_lda:

    mus = pm.Normal('mus', mu=0, sd=10, shape=2)
    sigma = pm.HalfCauchy('sigma', 5)

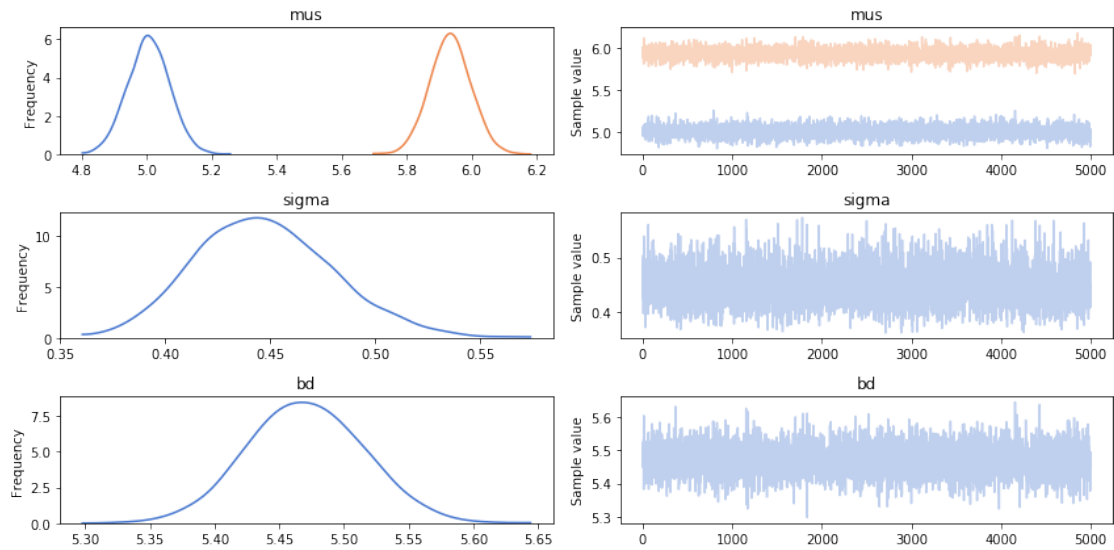
    setosa = pm.Normal('setosa', mu=mus[0], sd=sigma, observed=x_0[:50])
    versicolor = pm.Normal('versicolor', mu=mus[1], sd=sigma, observed=x_0[50:])

    bd = pm.Deterministic('bd', (mus[0]+mus[1])/2)

    trace_lda = pm.sample(5000, chains=1, njobs=1)

Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Sequential sampling (1 chains in 1 job)
NUTS: [sigma, mus]
100%|| 5500/5500 [00:03<00:00, 1703.58it/s]
Only one chain was sampled, this makes it impossible to run some convergence checks

In [36]: pm.traceplot(trace_lda);
```





```
In [37]: pm.summary(trace_lda)
```

```
Out[37]:
```

	mean	sd	mc_error	hpd_2.5	hpd_97.5
<code>mus__0</code>	5.01	0.06	9.93e-04	4.88	5.13
<code>mus__1</code>	5.93	0.06	9.24e-04	5.81	6.06
<code>sigma</code>	0.45	0.03	4.35e-04	0.38	0.51
<code>bd</code>	5.47	0.04	6.29e-04	5.39	5.56

```
In [38]: plt.axvline(trace_lda['bd'].mean(), ymax=1, color='r');
          bd_hpd = pm.hpd(trace_lda['bd']);
          plt.fill_betweenx([0, 1], bd_hpd[0], bd_hpd[1], color='r', alpha=0.5);
          plt.plot(x_0, y_0, 'o', color='k');
          plt.xlabel(x_n[0], fontsize=16);
```



```
In [39]: import sys, IPython, scipy, matplotlib, platform
         print("This notebook was created on a %s computer running %s and using:\nPython %s\nIP
```

This notebook was created on a x86_64 computer running debian buster/sid and using:

```
Python 3.7.2
IPython 7.2.0
PyMC3 3.6
NumPy 1.16.0
SciPy 1.2.0
Matplotlib 3.0.2
Seaborn 0.9.0
Pandas 0.23.4
```

```
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/ipykernel_launcher.py:2: DeprecationWarning: The 'warn' argument is deprecated in favor of the 'warns' argument, which is a list of warnings to raise. See https://docs.python.org/3/library/warnings.html#the-warn-function for more details.
```