

# Bayesian Data Analysis Chapter 6

January 20, 2019

```
In [1]: %matplotlib inline
import pymc3 as pm
import numpy as np
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib
palette = 'muted'
sns.set_palette(palette); sns.set_color_codes(palette)
```

## 1 Overfitting vs underfitting


```
In [2]: x = np.array([4.,5.,6.,9.,12, 14.])
y = np.array([4.2, 6., 6., 9., 10, 10.])

order = [0, 1, 2, 5]
plt.plot(x, y, 'o')
for i in order:
    x_n = np.linspace(x.min(), x.max(), 100)
    coeffs = np.polyfit(x, y, deg=i)
    ffit = np.polyval(coeffs, x_n)

    p = np.poly1d(coeffs)
    yhat = p(x)
    ybar = np.mean(y)
    ssreg = np.sum((yhat-ybar)**2)
    sstot = np.sum((y - ybar)**2)
    r2 = ssreg / sstot


    plt.plot(x_n, ffit, label='order {},  $R^2 = {:.2f}$ '.format(i, r2))

plt.legend(loc=2, fontsize=14);
plt.xlabel('$x$', fontsize=14);
plt.ylabel('$y$', fontsize=14, rotation=0);
fig = matplotlib.pyplot.gcf()
fig.set_size_inches(16,10)
```



```
In [3]: plt.figure(figsize=(8, 6))
x_values = np.linspace(-10, 10, 300)
for df in [1, 2, 5, 15]:
    distri = stats.laplace(scale=df)
    x_pdf = distri.pdf(x_values)
    plt.plot(x_values, x_pdf, label='Laplace with scale ($b$) = {}'.format(df))


x_pdf = stats.norm.pdf(x_values)
plt.plot(x_values, x_pdf, label='Gaussian')
plt.xlabel('x')
plt.ylabel('p(x)', rotation=0)
plt.legend(loc=0, fontsize=14)
plt.xlim(-7, 7);
fig = matplotlib.pyplot.gcf()
fig.set_size_inches(16,10)
```



```
In [4]: x_1 = np.array([ 10.,  8., 13.,  9., 11., 14.,  6.,  4., 12.,  7.,  5.])
        y_1 = np.array([ 8.04,  6.95,  7.58,  8.81,  8.33,  9.96,  7.24,  4.26,
                        10.84,  4.82,  5.68])
```

```
In [5]: np.random.seed(1)
        real_alpha = 4.25
        real_beta = [8.7, -1.2]
        data_size = 20
        noise = np.random.normal(0, 2, size=data_size)
        x_1 = np.linspace(0, 5, data_size)
        y_1 = real_alpha + real_beta[0] * x_1 + real_beta[1] * x_1**2 + noise
```

```
In [6]: order = 2#5
        x_1p = np.vstack([x_1**i for i in range(1, order+1)])
        x_1s = (x_1p - x_1p.mean(axis=1, keepdims=True))/x_1p.std(axis=1, keepdims=True)
        y_1s = (y_1 - y_1.mean())/y_1.std()
        plt.scatter(x_1s[0], y_1s);
        plt.xlabel('$x$', fontsize=16);
        plt.ylabel('$y$', fontsize=16, rotation=0);
```



### model\_1: Not including polynomial term

```
In [7]: with pm.Model() as model_1:
        alpha = pm.Normal('alpha', mu=0, sd=10)
        beta = pm.Normal('beta', mu=0, sd=1)
        epsilon = pm.HalfCauchy('epsilon', 5)

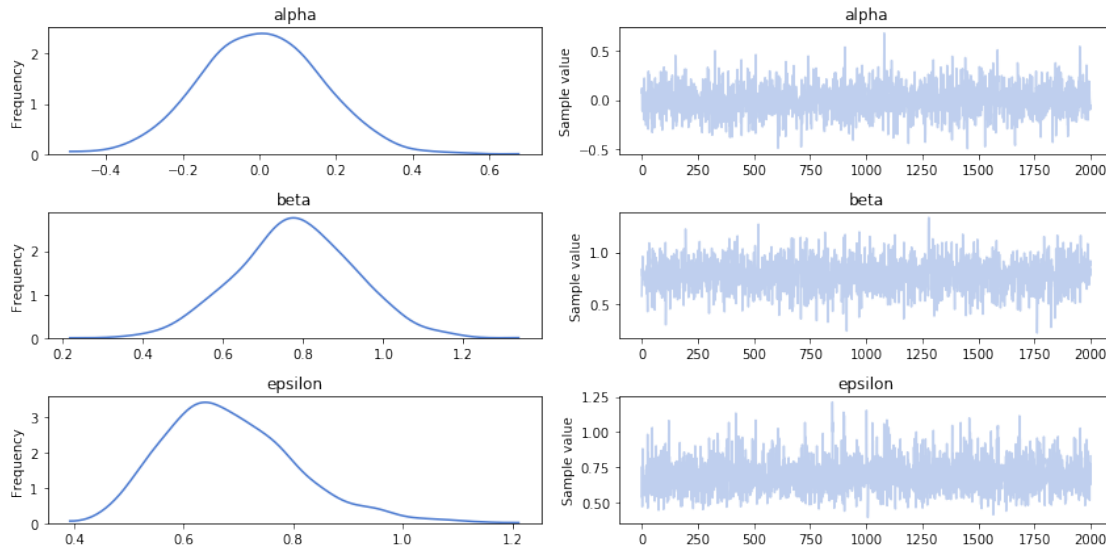
        mu = alpha + beta * x_1s[0]

        y_1 = pm.Normal('y_1', mu=mu, sd=epsilon, observed=y_1s)

        trace_1 = pm.sample(2100,chains=1,njobs=1)

Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Sequential sampling (1 chains in 1 job)
NUTS: [epsilon, beta, alpha]
100%|| 2600/2600 [00:02<00:00, 1189.52it/s]
Only one chain was sampled, this makes it impossible to run some convergence checks
```

```
In [8]: chain_1 = trace_1[100:]
        pm.traceplot(chain_1);
```



```
In [9]: pm.summary(chain_1)
```

```
Out[9]:
```

	mean	sd	mc_error	hpd_2.5	hpd_97.5
alpha	0.002707	0.157234	0.004058	-0.297762	0.309386
beta	0.779145	0.148678	0.003046	0.482230	1.051304
epsilon	0.683799	0.120173	0.002663	0.459332	0.920043

### model\_p: Including polynomial term

```
In [10]: with pm.Model() as model_p:
    alpha = pm.Normal('alpha', mu=0, sd=10)
    beta = pm.Normal('beta', mu=0, sd=1, shape=x_1s.shape[0])
    epsilon = pm.HalfCauchy('epsilon', 5)

    mu = alpha + pm.math.dot(beta, x_1s)

    y_1 = pm.Normal('y_1', mu=mu, sd=epsilon, observed=y_1s)

    trace_p = pm.sample(2100, chains=1, njobs=1)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt\_diag...

Sequential sampling (1 chains in 1 job)

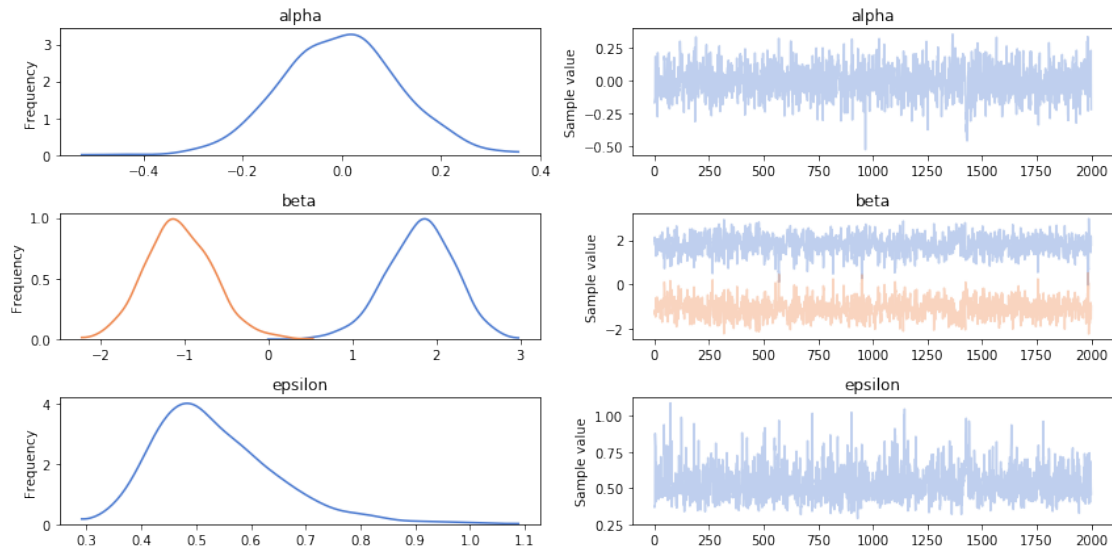
NUTS: [epsilon, beta, alpha]

100%|| 2600/2600 [00:04<00:00, 575.17it/s]

The acceptance probability does not match the target. It is 0.8964686391665999, but should be 0.9

Only one chain was sampled, this makes it impossible to run some convergence checks

```
In [11]: chain_p = trace_p[100:]
pm.traceplot(chain_p);
```



```
In [12]: pm.summary(chain_p)
```

```
Out[12]:
```

	mean	sd	mc_error	hpd_2.5	hpd_97.5
alpha	-0.006345	0.118836	0.003146	-0.220809	0.234278
beta__0	1.796237	0.404589	0.014456	0.996769	2.578225
beta__1	-1.061327	0.403376	0.014966	-1.888992	-0.312826
epsilon	0.535186	0.113336	0.003481	0.339629	0.762821


```
In [13]: alpha_l_post = chain_l['alpha'].mean()
betas_l_post = chain_l['beta'].mean(axis=0)
idx = np.argsort(x_1s[0])
y_l_post = alpha_l_post + betas_l_post * x_1s[0]

plt.plot(x_1s[0][idx], y_l_post[idx], label='Linear')

alpha_p_post = chain_p['alpha'].mean()
betas_p_post = chain_p['beta'].mean(axis=0)
y_p_post = alpha_p_post + np.dot(betas_p_post, x_1s)

plt.plot(x_1s[0][idx], y_p_post[idx], label='Pol order {}'.format(order))

plt.scatter(x_1s[0], y_1s)
plt.xlabel('$x$', fontsize=16)
plt.ylabel('$y$', fontsize=16, rotation=0);
plt.legend();
```



## 2 Information criteria

```
In [14]: model_p.name='Polynomial'
         model_l.name='Linear'
         dic_l = pm.compare({model_p: trace_p, model_l: trace_l})
         dic_l
```

```
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/pymc3/stats.py:167: FutureWarning
    return np.stack(logp)
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/pymc3/stats.py:167: FutureWarning
    return np.stack(logp)
```

```
Out[14]:
```

	WAIC	pWAIC	dWAIC	weight	SE	dSE	var_warn
Polynomial	32.67	2.64	0	1	4.61	0	0
Linear	42.5	2.14	9.83	0	4.16	2.71	0

```
In [15]: waic_l = pm.waic(trace=trace_l, model=model_l)
         waic_l
```

```
Out[15]: WAIC_r(WAIC=42.50361731898587, WAIC_se=4.157641613844861, p_WAIC=2.1412246201899223, v
```

```
In [16]: waic_p = pm.waic(trace=trace_p, model=model_p)
         waic_p
```

```

Out[16]: WAIC_r(WAIC=32.67190882656928, WAIC_se=4.610510226210436, p_WAIC=2.6390976063746674, v
In [17]: loo_l = pm.loo(trace=trace_l, model=model_l)
          loo_l

Out[17]: L00_r(L00=42.62390890178155, L00_se=4.1873713733381175, p_L00=2.2013704115877637, sha
In [18]: loo_p = pm.loo(trace=trace_p, model=model_p)
          loo_p

Out[18]: L00_r(L00=32.83617091380626, L00_se=4.654341727757128, p_L00=2.721228649993158, shape

```

### 2.0.1 Lower is better

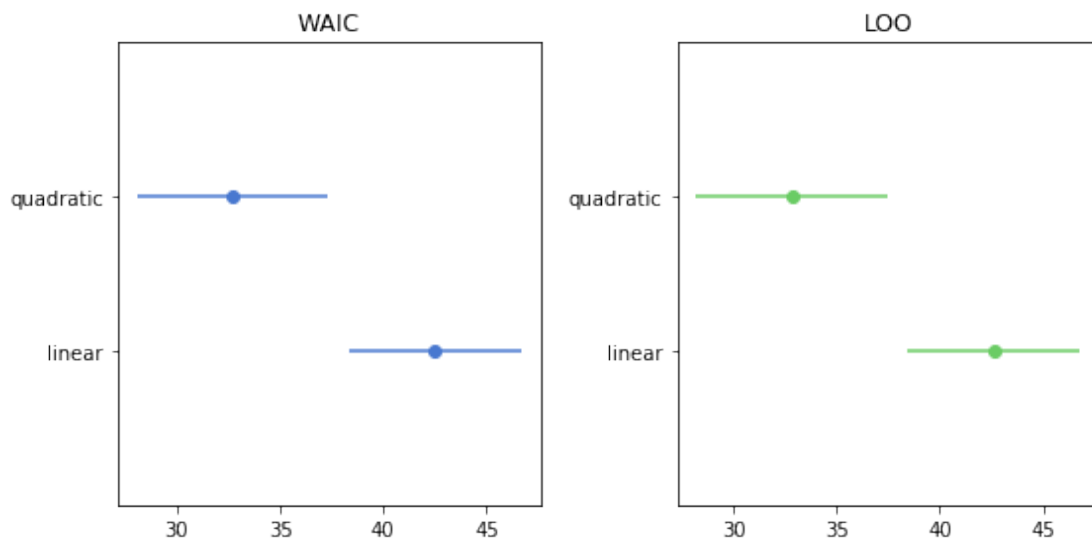
```

In [19]: plt.figure(figsize=(8, 4))
          plt.subplot(121)
          for idx, ic in enumerate((waic_l, waic_p)):
              plt.errorbar(ic[0], idx, xerr=ic[1], fmt='bo')
          plt.title('WAIC')
          plt.yticks([0, 1], ['linear', 'quadratic'])
          plt.ylim(-1, 2)

          plt.subplot(122)
          for idx, ic in enumerate((loo_l, loo_p)):
              plt.errorbar(ic[0], idx, xerr=ic[1], fmt='go')
          plt.title('L00')
          plt.yticks([0, 1], ['linear', 'quadratic'])
          plt.ylim(-1, 2)

          plt.tight_layout()

```






## 2.1 Posterior predictive checks

```
In [20]: plt.figure(figsize=(12,6))
plt.subplot(121)
plt.scatter(x_1s[0], y_1s, c='r');
plt.ylim(-3, 3)
plt.xlabel('x')
plt.ylabel('y', rotation=0)
plt.title('Linear')

for i in range(0, len(chain_1['alpha']), 50):
    plt.scatter(x_1s[0], chain_1['alpha'][i] + chain_1['beta'][i]*x_1s[0], c='g',
                edgecolors='g', alpha=0.5);
plt.plot(x_1s[0], chain_1['alpha'].mean() + chain_1['beta'].mean()*x_1s[0], c='g', al

plt.subplot(122)
plt.scatter(x_1s[0], y_1s, c='r');
plt.ylim(-3, 3)
plt.xlabel('x')
plt.ylabel('y', rotation=0)
plt.title('Order {}'.format(order))
for i in range(0, len(chain_p['alpha']), 50):
    plt.scatter(x_1s[0], chain_p['alpha'][i] + np.dot(chain_p['beta'][i], x_1s), c='g',
                edgecolors='g', alpha=0.5)
idx = np.argsort(x_1)
plt.plot(x_1s[0][idx], alpha_p_post + np.dot(betas_p_post, x_1s)[idx], c='g', alpha=1.

plt.tight_layout()
```



### 3 Bayes factors

```
In [21]: coins = 30 # 300
        heads = 9 # 90
        y = np.repeat([0, 1], [coins-heads, heads])
        print('These are the coin tosses we are modelling:\n {}'.format(y))
```

These are the coin tosses we are modelling:

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1]
```

```
In [22]: with pm.Model() as model_BF:
        p = np.array([0.5, 0.5])
        # model_index is a stochastic variable governed by the Categorical distribution
        # returning 0 or 1 for each model respectively
        model_index = pm.Categorical('model_index', p=p)
        # there are two models with different priors
        # one alpha=4,beta=8 and another alpha=8,beta=4
        m_0 = (4, 8)
        m_1 = (8, 4)
        # m returns the alpha,betas based on whether model index ==0
        # or not (`pm.math.eq(model_index, 0)`)
        m = pm.math.switch(pm.math.eq(model_index, 0), m_0, m_1)
        # prior on theta of the Bernoulli
        theta = pm.Beta('theta', m[0], m[1])
        # likelihood, y-> 1 heads, 0 tails
        y_pred = pm.Bernoulli('y', theta, observed=y)

        trace_BF = pm.sample(5000,chains=1,njobs=1)
```

Sequential sampling (1 chains in 1 job)

CompoundStep

>BinaryGibbsMetropolis: [model\_index]

>NUTS: [theta]

100%|| 5500/5500 [00:03<00:00, 1774.92it/s]

Only one chain was sampled, this makes it impossible to run some convergence checks

$$4 \quad BayesFactor = \frac{p(y|M_0)}{p(y|M_1)} = \frac{p(M_0|y) \times p(M_1)}{p(M_1|y) \times p(M_0)}$$

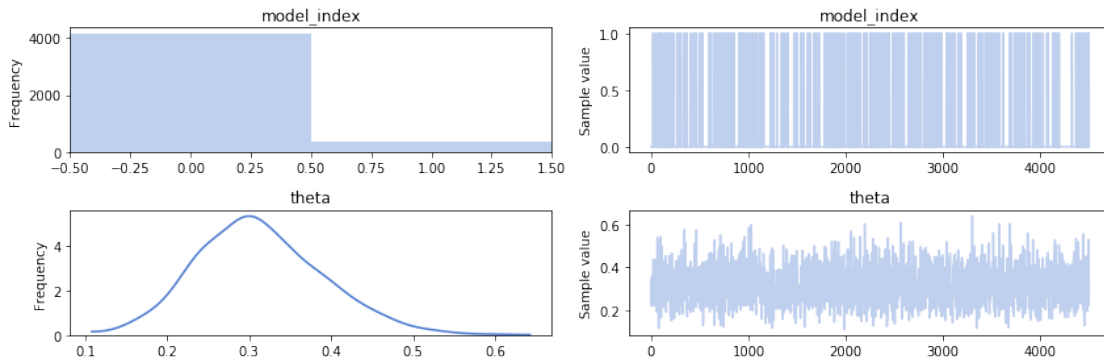
```
In [23]: chain_BF = trace_BF[500:]
        pm.traceplot(chain_BF)
        pM1 = chain_BF['model_index'].mean()
        pM0 = 1 - pM1
        print('Prior of selecting model 0 = {:.2f}'.format(p[0]))
```

```

print('Prior of selecting model 1 = {:.2f}'.format(p[1]))
print('Posterior mean of selecting model 0 = {:.2f}'.format(pM1))
print('Posterior mean of selecting model 1 = {:.2f}'.format(pM0))
print('Bayes factor = {:.2f}, thus model 0 is more likely'.format((pM0/pM1)*(p[1]/p[0]))

```

Prior of selecting model 0 = 0.50  
 Prior of selecting model 1 = 0.50  
 Posterior mean of selecting model 0 = 0.08  
 Posterior mean of selecting model 1 = 0.92  
 Bayes factor = 11.43, thus model 0 is more likely



**4.0.1** Also model 0  $\alpha = 4, \beta = 8$  is more compatible with the observations of mostly tails and  $\theta < 0.5$

## 4.1 Comparison of models using Information Criteria

```

In [24]: with pm.Model() as model_BF_0:
          theta = pm.Beta('theta', 4, 8)
          y_pred = pm.Bernoulli('y', theta, observed=y)

          trace_BF_0 = pm.sample(5000,chains=1,njobs=1)

```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt\_diag...

Sequential sampling (1 chains in 1 job)

NUTS: [theta]


100%|| 5500/5500 [00:02<00:00, 2580.73it/s]

Only one chain was sampled, this makes it impossible to run some convergence checks

```

In [25]: chain_BF_0 = trace_BF_0[500:]
          pm.traceplot(chain_BF_0);

```



```
In [26]: with pm.Model() as model_BF_1:
          theta = pm.Beta('theta', 8, 4)
          y_pred = pm.Bernoulli('y', theta, observed=y)

          trace_BF_1 = pm.sample(5000,chains=1,njobs=1)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt\_diag...


Sequential sampling (1 chains in 1 job)

NUTS: [theta]

100%|| 5500/5500 [00:02<00:00, 2729.76it/s]

Only one chain was sampled, this makes it impossible to run some convergence checks

```
In [27]: chain_BF_1 = trace_BF_1[500:]
          pm.traceplot(chain_BF_1);
```



```
In [28]: waic_0 = pm.waic(chain_BF_0, model_BF_0)
          waic_0
```

/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/pymc3/stats.py:167: FutureWarning  
return np.stack(logp)

Out[28]: WAIC\_r(WAIC=38.09485595294723, WAIC\_se=4.250210218039174, p\_WAIC=0.7153551590167374, v

```
In [29]: waic_1 = pm.waic(chain_BF_1, model_BF_1)
          waic_1
```

Out[29]: WAIC\_r(WAIC=39.45664158585026, WAIC\_se=2.0465734755680027, p\_WAIC=0.6998503578974625, v

```

In [30]: loo_0 = pm.loo(chain_BF_0, model_BF_0)
         loo_0

Out[30]: Loo_r(LOO=38.096185450169244, LOO_se=4.250485469459661, p_LOO=0.7160199076277465, shap

In [31]: loo_1 = pm.loo(chain_BF_1, model_BF_1)
         loo_1

Out[31]: Loo_r(LOO=39.45818767404322, LOO_se=2.0467180085626864, p_LOO=0.7006234019939441, shap

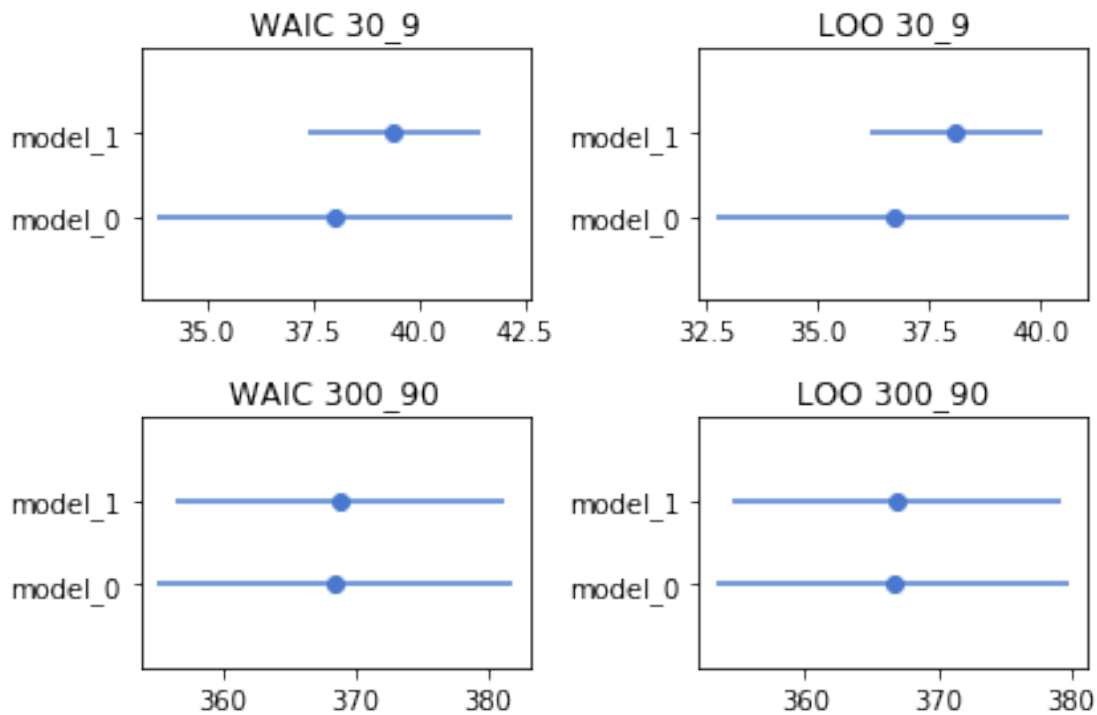
In [32]: est = [((38.02, 4.17), (39.41, 2.04)), ((36.69, 3.96), (38.09, 1.94)),
                ((368.41, 13.40), (368.76, 12.48)) , ((366.61, 13.31), (366.87, 12.34))]

         title = ['WAIC 30_9', 'LOO 30_9', 'WAIC 300_90', 'LOO 300_90']

         for i in range(4):
             plt.subplot(2,2,i+1)
             for idx, ic in enumerate(est[i]):
                 plt.errorbar(ic[0], idx, xerr=ic[1], fmt='bo')
             plt.title(title[i])
             plt.yticks([0, 1], ['model_0', 'model_1'])
             plt.ylim(-1, 2)

         plt.tight_layout()

```



**4.1.1** Bayes factors are sensitive to the selection of priors and this selection becomes less relevant as we increase the data. As you can see from using the Information Criteria there isn't much of a difference between the models, and these differences are going to become less and less as we increase the number of data. This sensitivity of Bayes factors to the selection of priors makes people not want to use them as much.

```
In [33]: import sys, IPython, scipy, matplotlib, platform
         print("This notebook was created on a %s computer running %s and using:\nPython %s\nIP
```

This notebook was created on a x86\_64 computer running debian buster/sid and using:

Python 3.7.2

IPython 7.2.0

PyMC3 3.6

NumPy 1.16.0

SciPy 1.2.0

Matplotlib 3.0.2

Seaborn 0.9.0

/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/ipykernel\_launcher.py:2: DeprecationWarning: The /path/to/ipykernel\_launcher.py module is deprecated. Use the ipykernel\_launcher.py module instead.