

# Bayesian Data Analysis Chapter 7

January 20, 2019

```
In [1]: %matplotlib inline
import pymc3 as pm
import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns
palette = 'muted'
sns.set_palette(palette); sns.set_color_codes(palette)
np.set_printoptions(precision=2)
```


## 1 Simple example

```
In [2]: clusters = 3

n_cluster = [90, 50, 75]
n_total = sum(n_cluster)

means = [9, 21, 35]
std_devs = [2, 2, 2]
# example of mixture data
mix = np.random.normal(np.repeat(means, n_cluster), np.repeat(std_devs, n_cluster))

In [3]: sns.kdeplot(np.array(mix))
plt.xlabel('$x$', fontsize=14);
```



In [4]: *# Author: Thomas Boggs*

```
import matplotlib.tri as tri
from functools import reduce
from matplotlib import ticker, cm

_corners = np.array([[0, 0], [1, 0], [0.5, 0.75**0.5]])
_triangle = tri.Triangulation(_corners[:, 0], _corners[:, 1])
_midpoints = [(_corners[(i + 1) % 3] + _corners[(i + 2) % 3]) / 2.0 for i in range(3)]

def xy2bc(xy, tol=1.e-3):
    '''Converts 2D Cartesian coordinates to barycentric.
    Arguments:
        xy: A length-2 sequence containing the x and y value.
    '''
    s = [(_corners[i] - _midpoints[i]).dot(xy - _midpoints[i]) / 0.75 for i in range(3)]
    return np.clip(s, tol, 1.0 - tol)

class Dirichlet(object):
    def __init__(self, alpha):
        '''Creates Dirichlet distribution with parameter `alpha`.'''
        from math import gamma
        from operator import mul
        self._alpha = np.array(alpha)
```

```

        self._coef = gamma(np.sum(self._alpha)) / reduce(mul, [gamma(a) for a in self._alpha])
def pdf(self, x):
    '''Returns pdf value for `x`.'''
    from operator import mul
    return self._coef * reduce(mul, [xx ** (aa - 1)
                                     for (xx, aa) in zip(x, self._alpha)])

def sample(self, N):
    '''Generates a random sample of size `N`.'''
    return np.random.dirichlet(self._alpha, N)


def draw_pdf_contours(dist, nlevels=100, subdiv=8, **kwargs):
    '''Draws pdf contours over an equilateral triangle (2-simplex).
    Arguments:
        dist: A distribution instance with a `pdf` method.
        border (bool): If True, the simplex border is drawn.
        nlevels (int): Number of contours to draw.
        subdiv (int): Number of recursive mesh subdivisions to create.
        kwargs: Keyword args passed on to `plt.triplot`.
    '''
    refiner = tri.UniformTriRefiner(_triangle)
    trimesh = refiner.refine_triangulation(subdiv=subdiv)
    pvals = [dist.pdf(xy2bc(xy)) for xy in zip(trimesh.x, trimesh.y)]

    plt.tricontourf(trimesh, pvals, nlevels, cmap=cm.Blues, **kwargs)
    plt.axis('equal')
    plt.xlim(0, 1)
    plt.ylim(0, 0.75**0.5)
    plt.axis('off')

In [5]: alphas = [[0.5] * 3, [1] * 3, [10] * 3, [2, 5, 10]]
for (i, alpha) in enumerate(alphas):
    plt.subplot(2, 2, i + 1)
    dist = Dirichlet(alpha)
    draw_pdf_contours(dist)

    plt.title(r'$\alpha$ = ({:.1f}, {:.1f}, {:.1f})'.format(*alpha), fontsize=16)

```



```
In [6]: with pm.Model() as model_kg:
        # Each observation is assigned to a cluster/component with probability p
        p = pm.Dirichlet('p', a=np.ones(clusters))
        category = pm.Categorical('category', p=p, shape=n_total)

        # Known Gaussians means
        means = pm.math.constant([10, 20, 35])


        y = pm.Normal('y', mu=means[category], sd=2, observed=mix)

        step1 = pm.ElemwiseCategorical(vars=[category], values=range(clusters))
        ## The CategoricalGibbsMetropolis is a recent addition to PyMC3
        ## I have not find the time yet to experiment with it.
        #step1 = pm.CategoricalGibbsMetropolis(vars=[category])
        step2 = pm.Metropolis(vars=[p])
        trace_kg = pm.sample(10000, step=[step1, step2], chains=1, njobs=1)
```

```
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/ipykernel_launcher.py:11: DeprecationWarning:
  # This is added back by InteractiveShellApp.init_path()
Sequential sampling (1 chains in 1 job)
CompoundStep
>ElemwiseCategorical: [category]
>Metropolis: [p]
100%|| 10500/10500 [00:30<00:00, 349.40it/s]
```

Only one chain was sampled, this makes it impossible to run some convergence checks

```
In [7]: chain_kg = trace_kg[1000:]
varnames_kg = ['p']
pm.traceplot(chain_kg, varnames_kg);
```



```
In [8]: print(pm.summary(chain_kg, varnames_kg))
n_cluster_array = np.array(n_cluster)
print('\n')
print('Actual values of cluster fractions: {}'.format(n_cluster_array/n_cluster_array..
```

	mean	sd	mc_error	hpd_2.5	hpd_97.5
p__0	0.418133	0.032734	0.001302	0.347043	0.475764
p__1	0.234513	0.029194	0.001281	0.173813	0.285907
p__2	0.347354	0.032026	0.001378	0.281533	0.406833

Actual values of cluster fractions: [0.42 0.23 0.35]

```
In [9]: with pm.Model() as model_ug:
# Each observation is assigned to a cluster/component with probability p
p = pm.Dirichlet('p', a=np.ones(clusters))
category = pm.Categorical('category', p=p, shape=n_total)

# We estimate the unknown gaussians means and standard deviation
means = pm.Normal('means', mu=[10, 20, 35], sd=2, shape=clusters)
sd = pm.HalfCauchy('sd', 5)

y = pm.Normal('y', mu=means[category], sd=sd, observed=mix)

step1 = pm.ElemwiseCategorical(vars=[category], values=range(clusters))
step2 = pm.Metropolis(vars=[means, sd, p])
trace_ug = pm.sample(10000, step=[step1, step2], chains=1, njobs=1)
```

```
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/ipykernel_launcher.py:12: DeprecationWarning:
if sys.path[0] == '':
Sequential sampling (1 chains in 1 job)
CompoundStep
```

```

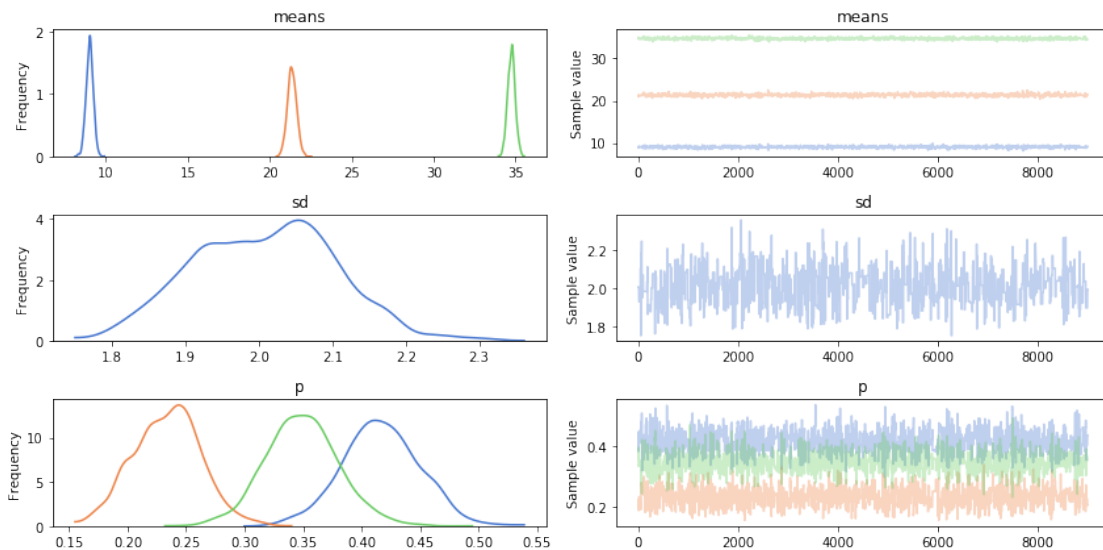
>ElemwiseCategorical: [category]
>CompoundStep
>>Metropolis: [p]
>>Metropolis: [sd]
>>Metropolis: [means]
100%|| 10500/10500 [00:35<00:00, 297.53it/s]
Only one chain was sampled, this makes it impossible to run some convergence checks

```

```

In [10]: chain_ug = trace_ug[1000:]
varnames_ug = ['means', 'sd', 'p']
pm.traceplot(chain_ug, varnames_ug);

```



```

In [11]: print(pm.summary(chain_ug, varnames_ug))
means_array = np.array([9, 21, 35])
std_devs = [2, 2, 2]
print('\n')
print('Actual values of cluster fractions: {}'.format(n_cluster_array/n_cluster_array))
print('Actual values of cluster means: {}'.format(means_array))
print('Actual values of cluster sd: {}'.format(std_devs))

```


	mean	sd	mc_error	hpd_2.5	hpd_97.5
means__0	9.008717	0.217561	0.007712	8.613627	9.463408
means__1	21.356761	0.275587	0.010145	20.812555	21.887742
means__2	34.742823	0.225639	0.008249	34.304095	35.202331
sd	2.008365	0.097757	0.004210	1.821559	2.183144
p__0	0.417117	0.033083	0.001022	0.351142	0.480680
p__1	0.233721	0.028664	0.000926	0.172529	0.283838
p__2	0.349161	0.030978	0.001001	0.288788	0.411532

```
Actual values of cluster fractions: [0.42 0.23 0.35]
Actual values of cluster means: [ 9 21 35]
Actual values of cluster sd: [2, 2, 2]
```

```
In [12]: ppc = pm.sample_ppc(chain_ug, 50, model_ug)
```

```
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/ipykernel_launcher.py:1: DeprecationWarning:
  """Entry point for launching an IPython kernel.
100%|| 50/50 [00:00<00:00, 998.86it/s]
```

```
In [13]: for i in ppc['y']:
          sns.kdeplot(i, alpha=0.1, color='b')
          sns.kdeplot(np.array(mix), lw=2, color='k');
          plt.xlabel('$x$', fontsize=14);
```



**Note the higher uncertainty where the data overlap and the reduced uncertainty at the high/low limits**

## 1.1 Marginalized Gaussian Mixture model

1.1.1 In the previous models we have explicitly defined the latent variable  $z$  in the model. This is ineffective in terms of sampling. PyMC3 offers the ability to model the outcome conditionally on  $z$  as  $p(y|z, \theta)$  and marginalise it to get  $p(y|\theta)$

```
In [14]: with pm.Model() as model_mg:
          p = pm.Dirichlet('p', a=np.ones(clusters))

          means = pm.Normal('means', mu=[10, 20, 35], sd=2, shape=clusters)
          sd = pm.HalfCauchy('sd', 5)

          y = pm.NormalMixture('y', w=p, mu=means, sd=sd, observed=mix)

          trace_mg = pm.sample(5000, chains=1, njobs=1)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt\_diag...

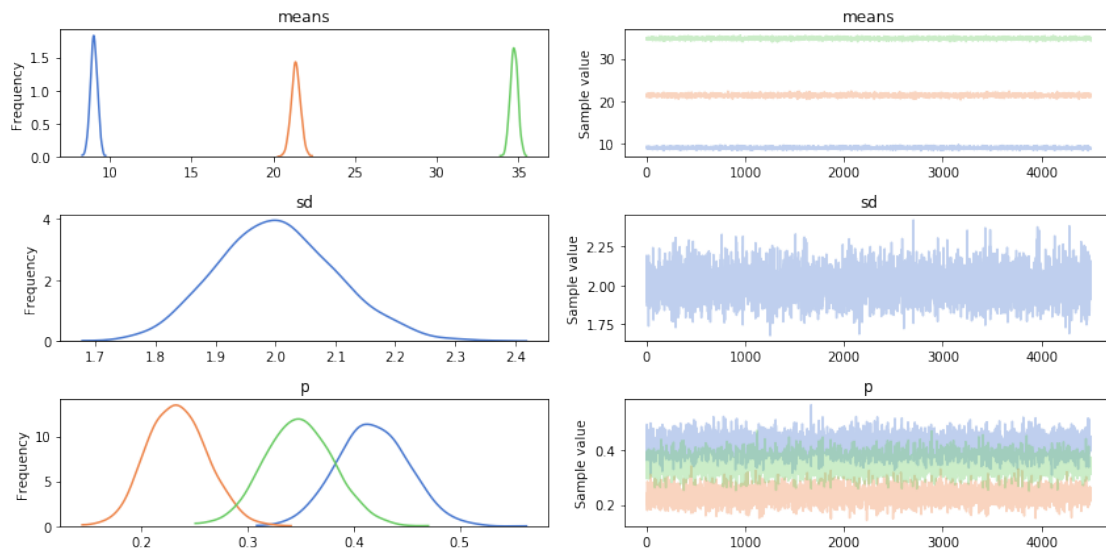
Sequential sampling (1 chains in 1 job)

NUTS: [sd, means, p]

100%|| 5500/5500 [00:08<00:00, 637.72it/s]

Only one chain was sampled, this makes it impossible to run some convergence checks


```
In [15]: chain_mg = trace_mg[500:]
          varnames_mg = ['means', 'sd', 'p']
          pm.traceplot(chain_mg, varnames_mg);
```





## 1.2 Zero inflated Poisson model


```
In [16]: lam_params = [0.5, 1.5, 3, 8]
k = np.arange(0, max(lam_params) * 3)
for lam in lam_params:
    y = stats.poisson(lam).pmf(k)
    plt.plot(k, y, 'o-', label="$\\lambda$ = {:.1f}".format(lam))
plt.legend();
plt.xlabel('$k$', fontsize=14);
plt.ylabel('$pmf(k)$', fontsize=14);
```



```
In [17]: np.random.seed(42)
n = 100
lam_true = 2.5 # Poisson rate
pi = 0.2 # probability of extra-zeros (pi = 1-psi)

# Simulate some data
counts = np.array([(np.random.random() > pi) * np.random.poisson(lam_true) for i in range(n)])

In [18]: plt.hist(counts, bins=30);
```

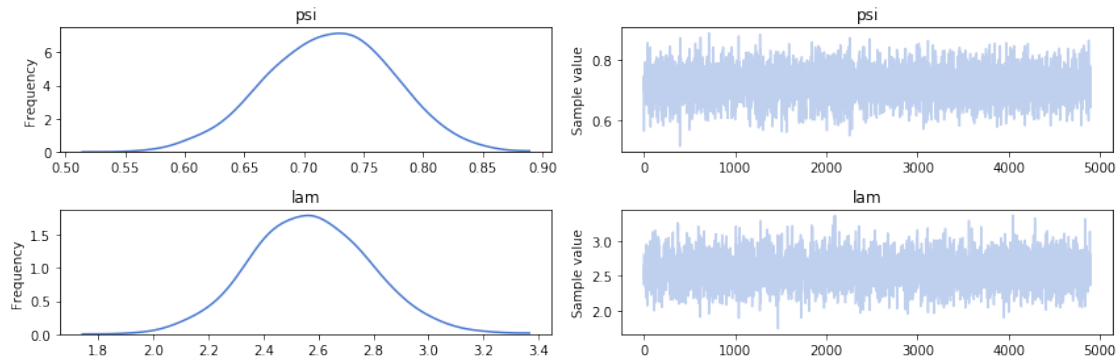


```
In [19]: with pm.Model() as ZIP:
          psi = pm.Beta('psi', 1, 1)
          lam = pm.Gamma('lam', 2, 0.1)

          y_pred = pm.ZeroInflatedPoisson('y_pred', theta=lam, psi=psi, observed=counts)
          trace_ZIP = pm.sample(5000, chains=1, njobs=1)
```

```
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Sequential sampling (1 chains in 1 job)
NUTS: [lam, psi]
100%|| 5500/5500 [00:03<00:00, 1652.34it/s]
Only one chain was sampled, this makes it impossible to run some convergence checks
```

```
In [20]: chain_ZIP = trace_ZIP[100:]
          pm.traceplot(chain_ZIP);
```



```
In [21]: pm.summary(chain_ZIP)
```

```
Out[21]:
```

	mean	sd	mc_error	hpd_2.5	hpd_97.5
psi	0.721471	0.052295	0.000769	0.619399	0.823200
lam	2.572429	0.214952	0.003234	2.154592	3.001078

### 1.3 Zero inflated Poisson regression

```
In [22]: #Kruschke plot
```


```
In [23]: fish_data = pd.read_csv('fish.csv')
fish_data.head()
```

```
Out[23]:
```

	nofish	livebait	camper	persons	child	xb	zg	count
0	1	0	0	1	0	-0.896315	3.050405	0
1	0	1	1	1	0	-0.558345	1.746149	0
2	0	1	0	1	0	-0.401731	0.279939	0
3	0	1	1	2	1	-0.956298	-0.601526	0
4	0	1	0	1	0	0.436891	0.527709	1

```
In [24]: plt.hist(fish_data['count'], bins=20, normed=True);
```

```
/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/matplotlib/axes/_axes.py:6521:
The 'normed' kwarg was deprecated in Matplotlib 2.1 and will be removed in 3.1. Use 'density' :
alternative="density", removal="3.1")
```



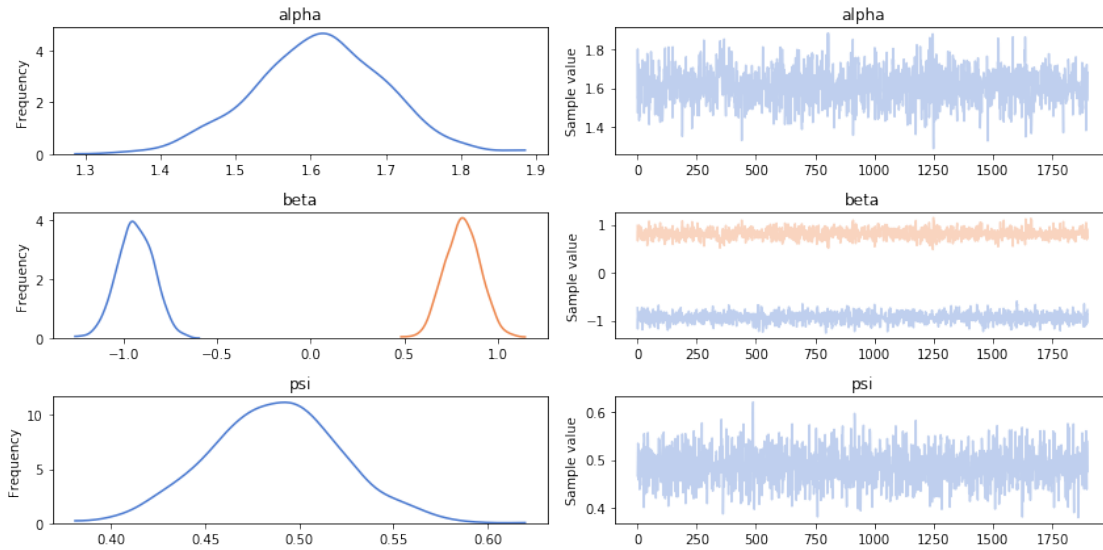
```
In [25]: with pm.Model() as ZIP_reg:
          psi = pm.Beta('psi', 1, 1)

          alpha = pm.Normal('alpha', 0, 10)
          beta = pm.Normal('beta', 0, 10, shape=2)
          lam = pm.math.exp(alpha + beta[0] * fish_data['child'] + beta[1] * fish_data['camp'])

          y = pm.ZeroInflatedPoisson('y', theta=lam, psi=psi, observed=fish_data['count'])
          trace_ZIP_reg = pm.sample(2000,chains=1,njobs=1)
```

```
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Sequential sampling (1 chains in 1 job)
NUTS: [beta, alpha, psi]
100%|| 2500/2500 [00:05<00:00, 437.66it/s]
Only one chain was sampled, this makes it impossible to run some convergence checks
```

```
In [26]: chain_ZIP_reg = trace_ZIP_reg[100:]
          pm.traceplot(chain_ZIP_reg);
```



```
In [27]: pm.summary(chain_ZIP_reg)
```

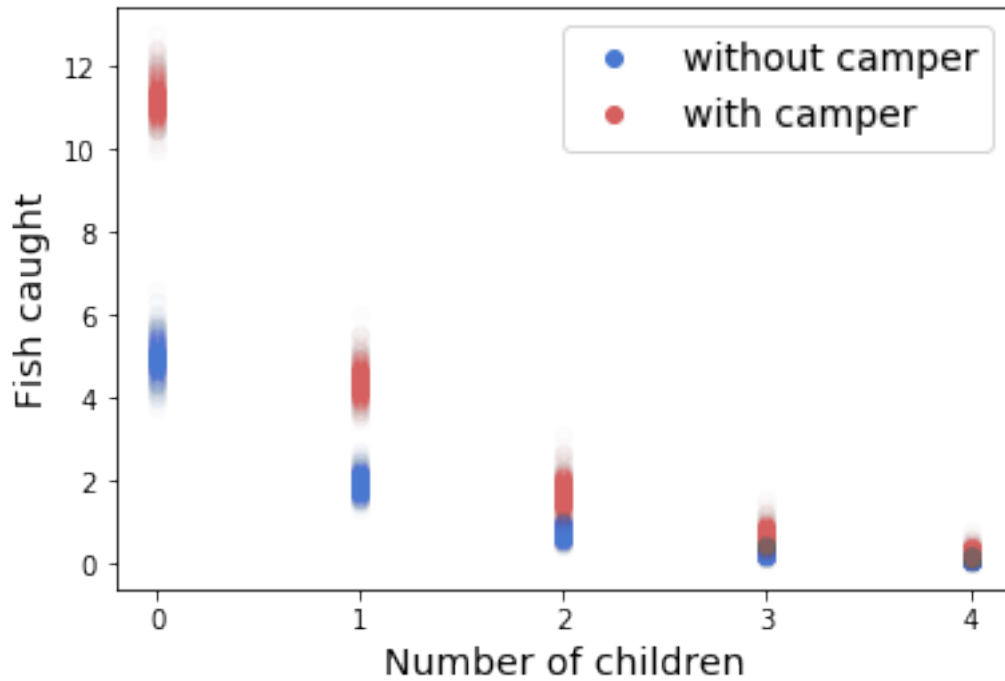
```
Out[27]:
```

	mean	sd	mc_error	hpd_2.5	hpd_97.5
alpha	1.610313	0.088094	0.002596	1.427976	1.771882
beta__0	-0.936188	0.097573	0.002547	-1.131112	-0.753889
beta__1	0.810109	0.096980	0.002897	0.621544	0.992202
psi	0.486522	0.034692	0.000816	0.419652	0.556741

```
In [28]: children = [0, 1, 2, 3, 4]
fish_count_pred_0 = []
fish_count_pred_1 = []
thin = 5
# calculate the expectation lambda with and without a camper for difference number of
# note lambda from the model is exp(a+bX)
for n in children:
    without_camper = chain_ZIP_reg['alpha'][:,thin] + chain_ZIP_reg['beta'][:,0][::thin]
    with_camper = without_camper + chain_ZIP_reg['beta'][:,1][::thin]
    fish_count_pred_0.append(np.exp(without_camper))
    fish_count_pred_1.append(np.exp(with_camper))
```


```
In [29]: plt.plot(children, fish_count_pred_0, 'bo', alpha=0.01)
plt.plot(children, fish_count_pred_1, 'ro', alpha=0.01)
```

```
plt.xticks(children);
plt.xlabel('Number of children', fontsize=14)
plt.ylabel('Fish caught', fontsize=14)
plt.plot([], 'bo', label='without camper')
plt.plot([], 'ro', label='with camper')
plt.legend(fontsize=14);
```



## 1.4 Robust logistic Regression

```
In [30]: iris = sns.load_dataset("iris")
df = iris.query("species == ('setosa', 'versicolor')")
y_0 = pd.Categorical(df['species']).codes
x_n = 'sepal_length'
x_0 = df[x_n].values
# contaminate our set with ones with unusually small sepal length
y_0 = np.concatenate((y_0, np.ones(6)))
x_0 = np.concatenate((x_0, [4.2, 4.5, 4.0, 4.3, 4.2, 4.4]))
x_0_m = x_0 - x_0.mean()
plt.plot(x_0, y_0, 'o', color='k');
```



```
In [31]: with pm.Model() as model_rlg:
          alpha_tmp = pm.Normal('alpha_tmp', mu=0, sd=100)
          beta = pm.Normal('beta', mu=0, sd=10)

          mu = alpha_tmp + beta * x_0_m
          theta = pm.Deterministic('theta', 1 / (1 + pm.math.exp(-mu)))
          # add the mixture here as a combination of the logistic derived theta
          # and a random pi from a Beta distribution
          pi = pm.Beta('pi', 1, 1)
          p = pi * 0.5 + (1 - pi) * theta
          # correct alpha from centering
          alpha = pm.Deterministic('alpha', alpha_tmp - beta * x_0.mean())
          bd = pm.Deterministic('bd', -alpha/beta)

          y1 = pm.Bernoulli('y1', p=p, observed=y_0)

          trace_rlg = pm.sample(2000, njobs=1,chains=1)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt\_diag...

Sequential sampling (1 chains in 1 job)

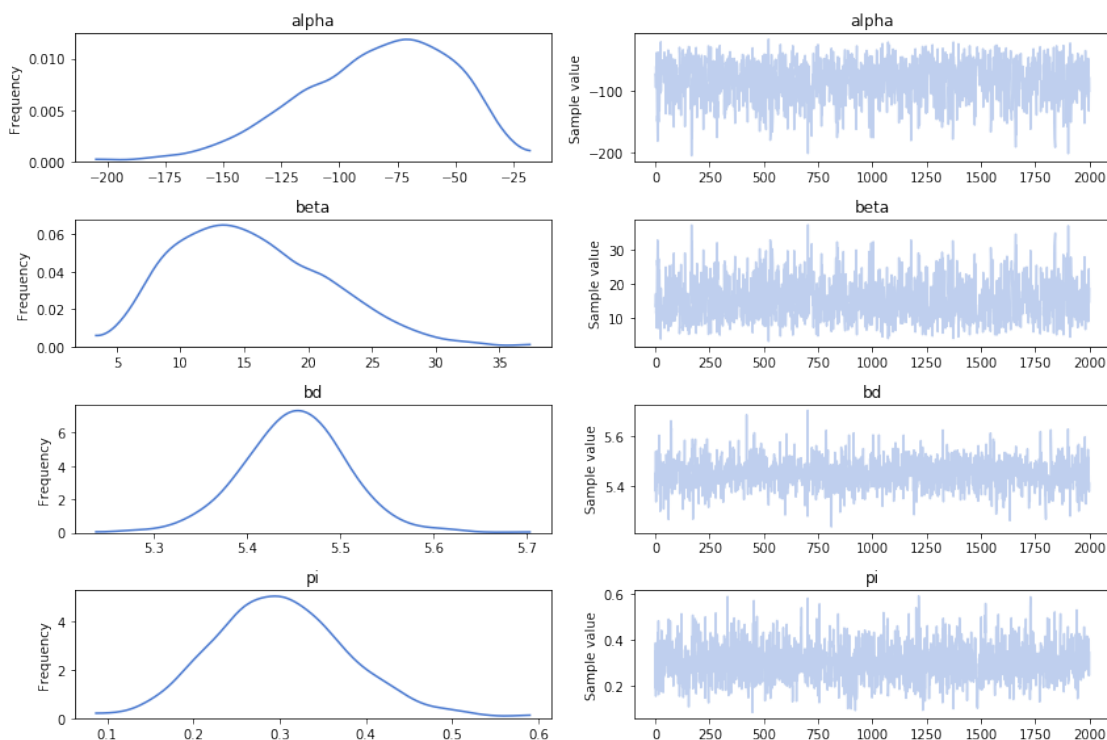
NUTS: [pi, beta, alpha\_tmp]

100%|| 2500/2500 [00:02<00:00, 1131.39it/s]

There were 3 divergences after tuning. Increase `target\_accept` or reparameterize.

Only one chain was sampled, this makes it impossible to run some convergence checks

```
In [32]: varnames = ['alpha', 'beta', 'bd', 'pi']
pm.traceplot(trace_rlg, varnames);
```



```
In [33]: pm.summary(trace_rlg, varnames)
```

```
Out[33]:
```


	mean	sd	mc_error	hpd_2.5	hpd_97.5
alpha	-84.752824	32.430609	1.204960	-145.679729	-28.152112
beta	15.543885	5.935414	0.219442	5.178316	26.580590
bd	5.450754	0.055642	0.001486	5.337383	5.553501
pi	0.300910	0.079280	0.002478	0.145793	0.453829

```
In [34]: theta = trace_rlg['theta'].mean(axis=0)
idx = np.argsort(x_0)
plt.plot(x_0[idx], theta[idx], color='b', lw=3);
plt.axvline(trace_rlg['bd'].mean(), ymax=1, color='r')
bd_hpd = pm.hpd(trace_rlg['bd'])
plt.fill_betweenx([0, 1], bd_hpd[0], bd_hpd[1], color='r', alpha=0.5)

plt.plot(x_0, y_0, 'o', color='k');
theta_hpd = pm.hpd(trace_rlg['theta'])[idx]
plt.fill_between(x_0[idx], theta_hpd[:,0], theta_hpd[:,1], color='b', alpha=0.5);

plt.xlabel(x_n, fontsize=16);
plt.ylabel('$\\theta$', rotation=0, fontsize=16);
```





```
In [35]: import sys, IPython, scipy, matplotlib, platform
         print("This notebook was created on a %s computer running %s and using:\nPython %s\nIP
```

This notebook was created on a x86\_64 computer running debian buster/sid and using:

Python 3.7.2

IPython 7.2.0

PyMC3 3.6

NumPy 1.16.0

SciPy 1.2.0

Matplotlib 3.0.2

Seaborn 0.9.0

Pandas 0.23.4

/home/damianos/miniconda3/envs/pymc3/lib/python3.7/site-packages/ipykernel\_launcher.py:2: DeprecationWarning: The 'warn' argument is deprecated since version 3.0. Use 'warn\_kwarg='warn\_kwarg' instead.